



# **Proposal for Noc benchmarking methodology**

Erno Salminen, Ari Kulmala

DACI research group

Tampere University of Technology

02.01.2007



# Noc WG goals (from white paper)

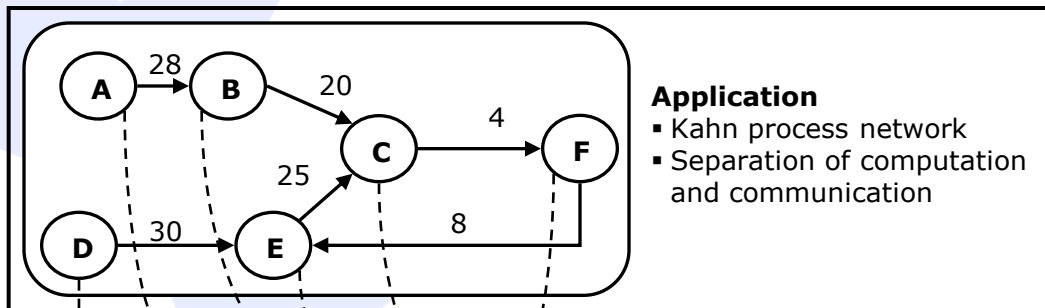
- Benchmarking environment should exercise
  - network size (small, medium, large)
  - IP core composition (amount of processing, memory cores, other)
  - Topology (regular, irregular)
  - Traffic characteristics (spatial and temporal)
  - QoS requirements
- Benchmark types
  - programs or program models (e.g. H.263)
  - micro-benchmarks (e.g. uniform traffic)

# Proposed methodology


- Separate
  - a) application functionality
  - b) mapping of application onto resources
  - c) hardware resources
  - d) benchmarked Noc
- Application model is fixed in benchmark suite
  - Includes dependencies
- Architecture and mapping may be modified during measurements in some cases

# Proposed methodology (2)


a)




b)



c)



d)



## Legend:

- Process (or task)
- Processing element (PE)

- Inter-process transfer
- - - Mapping of processes to PEs

- Application model (a) is interpreted by traffic generator (c) that models the resources
- Not all transfers directed to Noc
- E.g. transfer
  - A → B via Noc
  - E → C internal to PE3
  - They have different cost

# Application model

- Model similar to Kahn Process Network (KPN)
  - Vertices represent computation tasks
  - Edges represent communication channels
  - Channels carry the data tokens between tasks
- Model includes dependencies, timing, destination, and size of data transfers
- No actual computation, only external behavior of application tasks
  - Unlike approaches that include also computation in KPN (e.g. Artemis by Pimentel et al.)
- The task is triggered according to a *condition* that
  - depends on the received data tokens
  - and internal state of the task

# Application model (2)

- Internal state of the task is expressed with
  - an *execution counter*
  - and a *status variable*: {RUN, READY, WAIT, FREE}.
- The *dependence* of the multiple inputs can be either
  - a) OR: triggered when any of the input receives a token
  - b) AND: triggered when there is data in all inputs
    - Note: trad. KPN uses only AND dependency
- A *function* of a task defines the
  - *execution time*      in cycles
  - *output port(s)*      (their ID numbers)
  - *data amount*          to send in bytes

# Application model (3)

- A trigger includes one or several functions
- The function that is evaluated depends on the value of the *execution counter*.
- The *execution time* and *data amount* is expressed with either
  - a) a statistical distribution (uniform or normal)
  - b) a polynomial function which depend on the received data amount (scalar value is subset of this)
- Events (timers) can be used for environment modeling
- Some tasks may model memory accesses instead of computation

# Possible task models

## ■ Statistical

- e.g. destinations:
  - send to task 5 with 10% prob, to task 4 with 90 % prob and so on
- most compact format
- Either set probability to 100% or neglect dependencies to avoid locking

## ■ Modulo (model loops)

- e.g. two destinations:
  - to task 5 at every 10th iteration (modulo)
  - to task 4 otherwise

## ■ Trace

- runtime and data amount can be different for each execution
- large file



# Resource model

- Very coarse-grain
- Modeled by traffic generator
- PE has a performance of "*X operations/cycle*"
  - Possibly also area, power, max. frequency etc.
  - For Noc benchmarking, all resources may be identical (e.g. 1 ops/cycle)
  - Heterogeneous models needed for design space exploration and SoC modeling
- PEs may have different frequencies
- Example
  - task *i* has 500 operations
  - PEs *A* and *B* can execute 1 and 2 ops/cycle
  - task *i* takes
    - 500 cycles when mapped to PE *A* ( $=500/1$ )
    - 250 cycles when mapped to PE *B* ( $=500/2$ )

# Proposed file format

- Based on XML (eXtensible Markup Language)
  - ASCII text with user-specified tags
    - Larger file size but more reader-friendly than binary format
  - Format rules defined with *schema* file
  - Free parsers available for C, Tcl, Python
  - Free validators also available
- Benchmark tools convert XML to their internal format if needed
- Application, architecture, and mapping may be separate or combined into single file

# Example of task

- Triggered when data at input port 2
- Condition is the same for all exec counts
- Always sends 384 bytes to outport 0
  - function:  $0x + 384$
- Complexity is always 13571 integer operations
  - function:  $0x + 13571$
- Process is not finished for ever, but can be executed again (state=ready)

```
<application>
  <task_graph>

    <task name="dct" id="0" class="general">
      <in_port id="2"/>
      <out_port id="0"/>
      <trigger>
        <in_port id="2"/>
        <exec_count>
          <op_count>
            <int_ops>
              <polynomial>
                <param value="13571" exp="0"/>
              </polynomial>
            </int_ops>
          </op_count>
        </exec_count>
        <send out_id="0" prob="1">
          <byte_amount>
            <polynomial>
              <param value="384" exp="0"/>
            </polynomial>
          </byte_amount>
        </send>
        <next_state value="READY"/>
      </trigger>
    </task>
```

# Example of task (2)

```
<task name="VLCDecoding" id="9" class="general">
  <in_port id="7"/>
  <in_port id="21"/>
  <out_port id="11"/>
  <out_port id="20"/>
  <trigger dependence_type="or">
    <in_port id="7"/>
    <in_port id="21"/>
  <exec_count min="0" max="328" mod_period="330">
    <op_count>
```

```
...
<task_connection src="6" dst="7"/>
<task_connection src="20" dst="21"/>
...
```

```
<event_list>
  <event out_port_id="8" amount="1" trigger_type="periodic"
    name="VideoInput" id="11" time="0.0285714285714" prob="1"/>
</event_list>
```

## ■ Two inports: 7 and 21

- 7 receives data from output 6 (of another task), 21 from port 20

## ■ Two outputs: 11 and 20

## ■ Event "VideoInput" executes periodically every 0.028 seconds, it writes 1 byte to output 8

- Also "one-shot" events possible

## ■ Number of events, tasks, ports, or connections is not restricted

- 32-bit integers allow 4e9 unique IDs

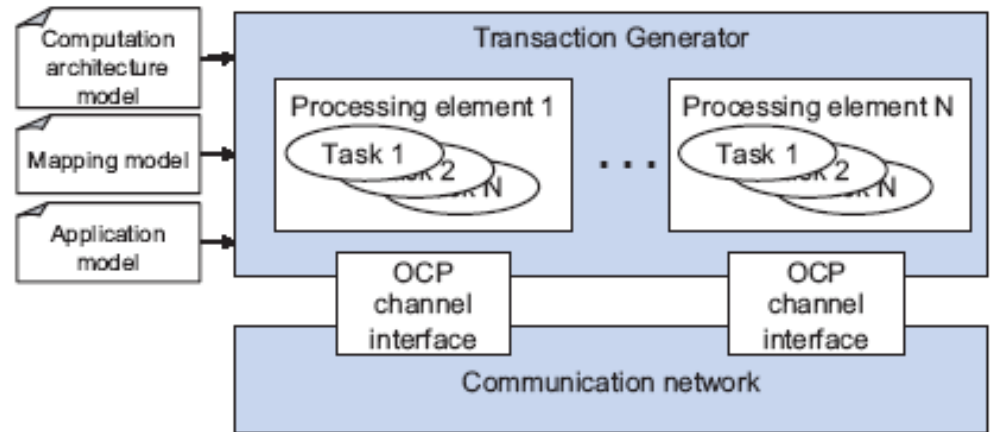


# Task mapping

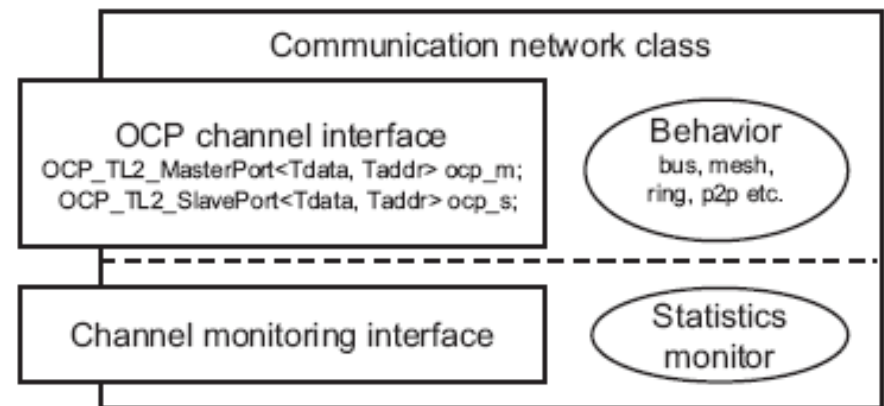
- Suitable mapping depends on
  - task and PE type
  - PE load (runtime)
  - induced communication
  - required memory vs available memory
- Tasks with (class=general) can be freely mapped to general purpose resources (CPUs)
- HW accelerator resource can execute only type of tasks
  - External interfaces (e.g. radio interface) are similar

# Transaction generator

- SystemC program that models the PEs
- Converts operation counts to time
- Sends, receives, and checks data
- Schedules the tasks mapped to same PE
  - Non-preemptive policy
  - FIFO, priority, round-robin



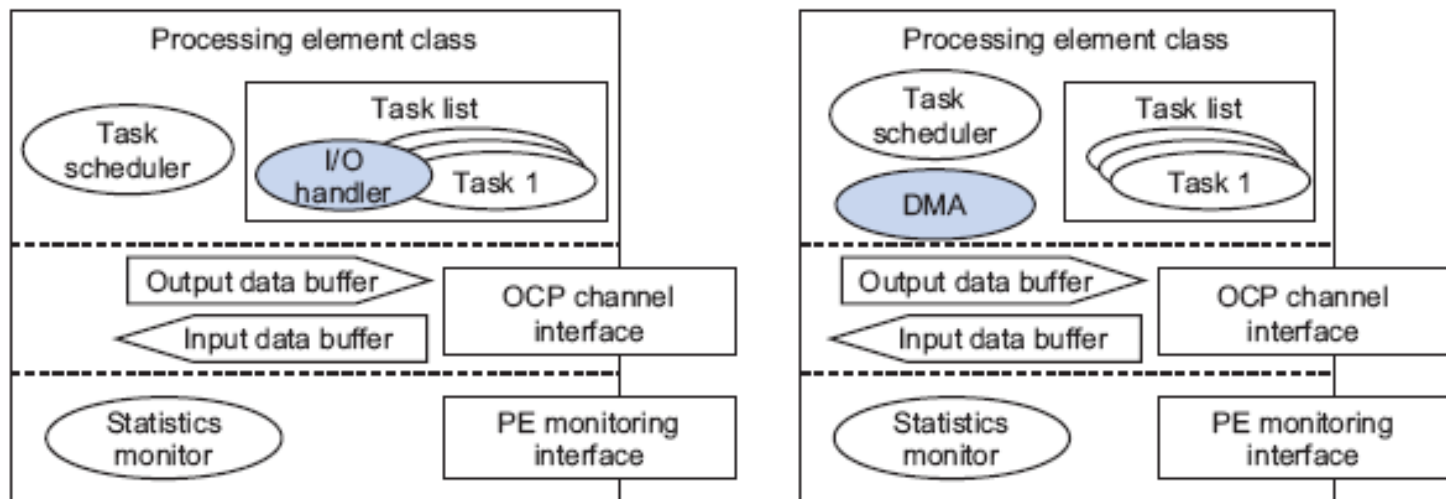
*Fig. 30. An overview of Transaction Generator implementation.*



*Fig. 31. Block diagram of the communication network class.*

# Modeling DMA

- Direct memory access allows
  - To send data to network while PE is processing
  - To receive data while processing
- Of course, DMA can be disabled
  - Processing not possible during the transfers




(a) PE class without DMA. The I/O handler runs sequentially with other tasks

(b) PE class with DMA. The DMA handles I/O transfers in parallel to other tasks

*Fig. 32. Processing element model for Transaction Generator with and without DMA modeling.*

# Task state machine

- After reset, state usually *Wait*
- Data token (tg\_packet) structure shown in Fig 34
- An internal transfer
  - communication between tasks located in the same PE
  - pass-by-reference (pointer)
- External transfers
  - carried out through an OCP TL2 interface
  - master port is utilized for outgoing, slave port for incoming transfers.



- a: Scheduler selected the task for execution  
b: The task has all the data for the next execution  
c: Fully executed but has not all data for the next execution  
d: Fully executed and is not dependent on any data anymore

*Fig. 33. Task execution scheduling state diagram.*

```
class tg_packet {  
    unsigned int _transfer_id;    // transfer id  
    unsigned int _source_id;     // source process id  
    unsigned int _target_id;     // target process id  
    unsigned int _n_bytes;       // amount of payload data  
    unsigned char* _data;        // pointer to payload  
}
```

*Fig. 34. The packet structure of a data transfer.*



# Statistics

- Some general statistics are collected automatically
- These are complemented with Noc-specific metrics




Fig. 35. Block diagram of the task class.

Table 9. Statistics produced by Transaction Generator.

Processing element	Task	Transfer
Task execution count	Execution count	Packet amount
Task execution time	Time in READY state	Min transfer latency
Read/write time	Time in RUN state	Avg transfer latency
Read/write packet count	Time in WAIT state	Max transfer latency
Read/write packet header amount	Time in FREE state	
Read/write packet payload amount	Last execution time	

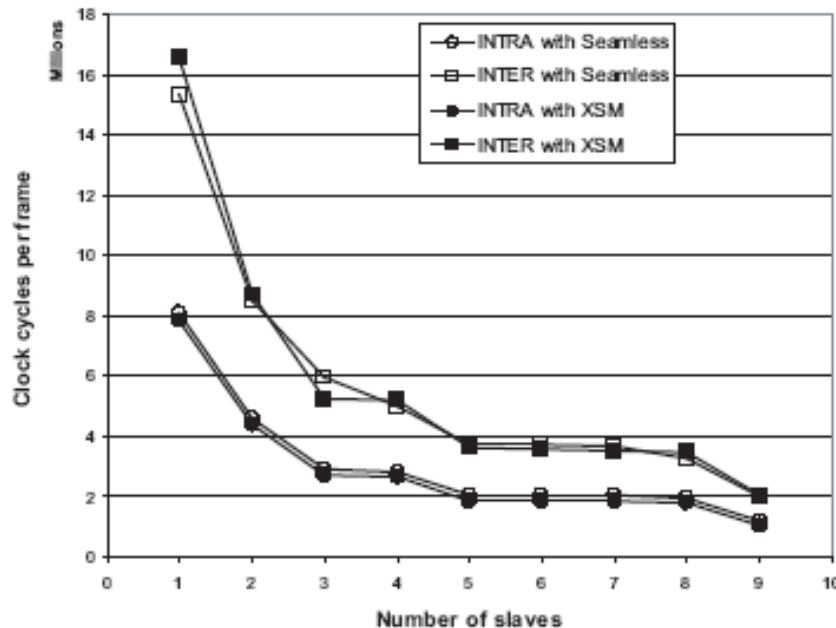
Additionally: required buffering

# Performance/real-time requirements

- Somewhat open subject currently
- E.g. define limits for certain paths
  - max time "task A starts executing" -> "task F has sent all the data to Noc" is 3.5 ms (= application runtime)
  - max time "A starts sending data" -> "B has received all the data from Noc" 0.6 ms
  - time limits in absolute time (1ms after reset this has happened, after 5 ms that has happened...)
  - These can be automatically checked
    - Define criticality level for requirements
      - a) Stop at violation
      - b) Continue despite
- If separate to application description
  - "Clean" application model
  - Problematic if data should be sent in different manner depending on requirements
  - Problematic if scheduler at the PE must react to requirements

# Accuracy

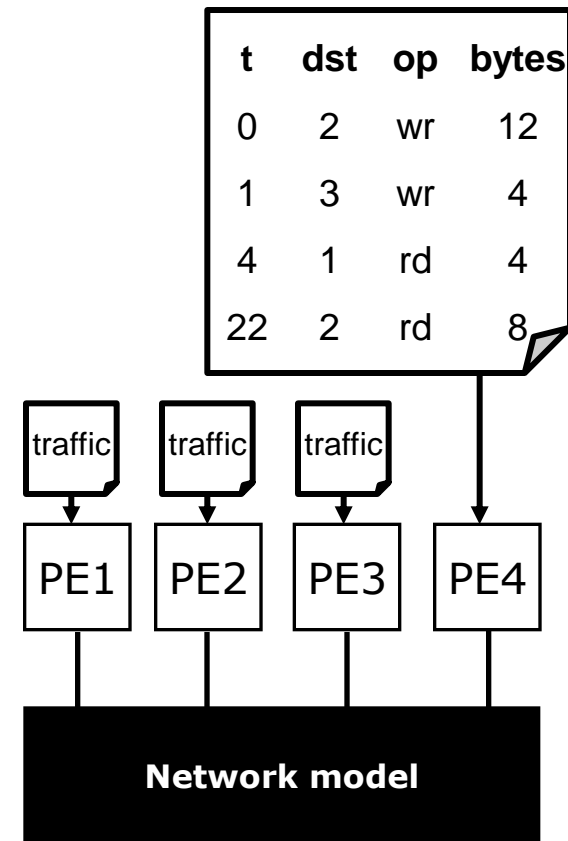
- Seems adequate for network comparison
- Not as issue with purely synthetic workloads



*Fig. 57. The clock cycle results obtained with XSM compared to ISS with Seamless for INTRA and INTER frames.*

# "Traditional" traffic model

- No dependencies
  - PEs might send results before receiving source data
- Traffic generated either
  - a) directly to Noc at runtime
  - b) to file at compile-time
  - The latter (b) allows exact reproduction easily
- PE type, task grouping and scheduling fixed
  - However, PE location within Noc may still be changed
- Not easy to determine PE load or application runtime
- Suitable for environment modeling (noise traffic, external IO)



# Benefits of proposed methodology

- Trad. case with one functional model per resource
  - Suitable for environment, HW accelerator modeling (e.g. DCT) and for micro-benchmarks
  - Not suitable for CPU modeling
    - multithreading and context switch overhead?
- Proposed method is a superset of trad. and allows
  - modeling dependencies between tasks
  - more tasks than resources: one resource may execute several tasks in time-multiplexed manner
  - CPU (+RTOS) modeling
  - real trace or statistical (e.g. avg. runtime) runtime
  - design space exploration (allocation, mapping, scheduling, Noc params)
  - also micro-benchmarks
    - tasks have self-loop, i.e. they trigger themselves repeatedly
    - same as neglecting dependencies

# Some related publications

- <http://www.tkt.cs.tut.fi/research/daci/> see "Publications" at bottom-left
- 1. Erno Salminen, Ari Kulmala, Timo D. Härmäläinen, "A brief history of Network-on-chip", Noc symposium (NOCS), Princeton, NJ, USA, May 7-9, 2007, submitted.
- 2. Erno Salminen, Tero Kangas, Vesa Lahtinen, Jouni Riihimäki, Kimmo Kuusilinna, Timo D. Härmäläinen, "Benchmarking Mesh and Hierarchical Bus Networks in System-on-Chip Context", Journal of System Architectures, 2007, Accepted.
- 3. Tero Kangas, "Methods and Implementations for Automated System on Chip Architecture Exploration", PhD Thesis, Tampere University of Technology, Publication 616, 2006, 181 pages.
- 4. Erno Salminen, Tero Kangas, Timo D. Härmäläinen, Jouni Riihimäki, "Requirements for Network-on-Chip Benchmarking", Norchip, Oulu, Finland, November 21-22, 2005, pp. 82-85.
- 5. Erno Salminen, Tero Kangas, Timo D. Härmäläinen, "The impact of communication on the scalability of the data-parallel video encoder on MPSoC", International Symposium on System-on-Chip, Tampere, Finland, November 14-16, 2006, pp. 191-194.
- 6. Ari Kulmala, Erno Salminen, Marko Hännikäinen, Timo D. Härmäläinen, "Evaluating SoC Network Performance in MPEG-4 Encoder", The IEEE 2006 Workshop on Signal Processing Systems (SiPS'06), Banff, Canada, October 2-4, 2006, pp. 271-276, IEEE.
- 7. Tero Kangas, Jouni Riihimäki, Erno Salminen, Kimmo Kuusilinna, Timo D. Härmäläinen, "Using a Communication Generator in SoC Architecture Exploration", International Symposium on System-on-Chip, Tampere, Finland, November 19-21, 2003, pp. 105-108.

